

Machine Learning 1

Tasnia Sharia (PID A15931128)

10/21/2021


First is clustering methods

Kmeans clustering

The function in base R to do Kmeans clustering is called 'kmeans()'

First make up some data where we know what the answer should be:

```
tmp <- c(rnorm(30,-3), rnorm(30,3))  
x <- cbind(x=tmp, y=rev(tmp))  
plot(x)
```



Q. Can we use `kmeans()` to cluster this data setting `k` to 2 and `nstart` to 20?

[illegible]

Q. How many points are in each cluster?

km\$size

```
## [1] 30 30
```

Q. What ‘component’ of your result object details cluster assignment/membership?

```
km$cluster
```

```
## [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1  
## [39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


Q. What ‘component’ of your result object details cluster center?

km\$centers

```
##          x          y
## 1  2.988375 -2.959826
## 2 -2.959826  2.988375
```

Q. Plot `x` colored by the `kmeans` cluster assignment and cluster centers as blue points

```
plot(x, col=km$cluster)
points(km$centers, col="blue", pch=15, cex=2)
```



Hierarchical Clustering

A big limitation with k-means is that we have to tell it K (the number of clusters we want).

Analyze this same data with `hclust()`


Demonstrate the use of `dist()`, `hclust()`, `plot()`, and `cutree()` function to do clustering. Generate dendeograms and return cluster assignment membership vector.

```
hc <- hclust( dist(x) )
hc

##
## Call:
## hclust(d = dist(x))
##
## Cluster method   : complete
## Distance         : euclidean
## Number of objects: 60
```

There is a plot method for `hclust` result objects. Let's see it.

```
plot(hc)
```



To get our cluster membership vector we have to do a wee bit more work. We have to “cut” the tree where we think it makes sense. For this we use the ‘`cutree()`’ function.

```
cutree(hc, h=8)
```


[illegible]

You can also call 'cutree()' setting k= the number of grps/clusters you want

```
grps <- cutree(hc, k=2)
```

Make our results plot

```
plot(x, col=grps)
```



Now, We will examine the PCA of UK food data

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
View(x)
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
nrow(x)
```

```
## [1] 17
```

```
ncol(x)
```

```
## [1] 5
```

```
dim(x)
```

```
## [1] 17 5
```

The `dim()` function returns the # of rows and columns. Using `nrow()` or `ncol()` functions also provide the answer

```
## Preview the first 6 rows
head(x)
```

```
##           X England Wales Scotland N.Ireland
## 1      Cheese      105   103      103       66
## 2 Carcass_meat     245   227      242      267
## 3   Other_meat     685   803      750      586
## 4        Fish     147   160      122       93
## 5 Fats_and_oils     193   235      184      209
## 6        Sugars     156   175      147      139
```

```
## Preview the last 6 rows
tail(x)
```

```
##           X England Wales Scotland N.Ireland
## 12  Fresh_fruit     1102  1137      957      674
## 13     Cereals     1472  1582     1462     1494
## 14   Beverages        57    73       53       47
## 15  Soft_drinks     1374  1256     1572     1506
## 16 Alcoholic_drinks    375   475      458      135
## 17 Confectionery       54    64       62       41
```

The data should be 17 by 4 dimensions. There is an extra first column that needs to be fixed.

```
# Note how the minus indexing works
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

```
##           England Wales Scotland N.Ireland
## Cheese      105   103      103       66
## Carcass_meat 245   227      242      267
## Other_meat   685   803      750      586
## Fish        147   160      122       93
## Fats_and_oils 193   235      184      209
## Sugars       156   175      147      139
```

Now let's check the dimensions again

```
dim(x)
```

```
## [1] 17  4
```

Alternative approach to setting the correct row-names

```
x <- read.csv(url, row.names=1)
head(x)
```

```
##           England Wales Scotland N.Ireland
## Cheese           105   103       103       66
## Carcass_meat     245   227       242       267
## Other_meat       685   803       750       586
## Fish             147   160       122        93
## Fats_and_oils    193   235       184       209
## Sugars           156   175       147       139
```

```
dim(x)
```

```
## [1] 17  4
```

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

The second approach to solving the ‘row-names problem’ where we include the `row.names` argument is more convenient.

```
# What would happened if we ran row.names(x) <- x[,1] again
rownames(x) <- x[,1]
head(x)
```

```
##           England Wales Scotland N.Ireland
## 105           105   103       103       66
## 245           245   227       242       267
## 685           685   803       750       586
## 147           147   160       122        93
## 193           193   235       184       209
## 156           156   175       147       139
```

By running `rownames(x) <- x[,1]` multiple times, the first column from the previous run would be removed. This would cause a problem of having data being removed after each run when the code is rewritten.

Now we go back to the original data with the right number of dimensions and correct headings.


```
x <- read.csv(url, row.names=1)
head(x)
```

```
##           England Wales Scotland N.Ireland
## Cheese           105   103       103       66
## Carcass_meat     245   227       242       267
## Other_meat       685   803       750       586
## Fish             147   160       122        93
## Fats_and_oils    193   235       184       209
## Sugars           156   175       147       139
```

```
View(x)
```


We will spot major differences and trends by generating boxplots and various pairwise plots that may not help as much

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above `barplot()` function results in the following plot?


```
barplot(as.matrix(x), beside=FALSE, col=rainbow(nrow(x)))
```

Changing the `beside` argument can change the plot. Setting `beside` argument to `FALSE` creates a stacked boxplot. Leaving the `beside` argument out would also give the same result because being `FALSE` is the default. If `beside=TRUE` were to be added, `TRUE` would be a numeric vector and includes other components that could be useful information to the data.

Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch=16)
```



It is difficult to make sense of this visual dataset. I cannot seem to make out what the figures represent. The axis are very confusing. The diagonal lines appear to show a trend but can't tell what sort of trend.

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

We cannot tell which figures belong to which countries. Does each country have its own respective row or column? Which figures belong to which country?

PCA to the rescue


`prcomp()` expects the observations to be rows and the variables to be columns therefore we need to first transpose our data.frame matrix with the `t()` transpose function.

```
# Use the prcomp() PCA function
pca <- prcomp( t(x) )
summary(pca)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4
## Standard deviation  324.1502 212.7478 73.87622 4.189e-14
## Proportion of Variance  0.6744  0.2905  0.03503 0.000e+00
## Cumulative Proportion  0.6744  0.9650  1.00000 1.000e+00
```


Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
# Plot PC1 vs PC2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x), col=c("orange", "red", "blue", "dark green"))
```



We automatically obtain information about the contributions of each PC to the total variance of the coordinates, which is contained in the Eigenvectors returned from such calculations

We can use the square of `pca$sdev`, which stands for “standard deviation”, to calculate how much variation in the original data each PC accounts for.

For the `prcomp()` function we can use the `summary()` command or examine the returned `pca$sdev` (see below).

```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```


```
## [1] 67 29 4 0
```

```
## summary() command for prcomp function
z <- summary(pca)
z$importance
```

```
##              PC1      PC2      PC3      PC4
## Standard deviation 324.15019 212.74780 73.87622 4.188568e-14
## Proportion of Variance 0.67444 0.29052 0.03503 0.000000e+00
## Cumulative Proportion 0.67444 0.96497 1.00000 1.000000e+00
```

This information can be summarized in a plot of the variances (eigenvalues) with respect to the principal component number (eigenvector number)


```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```



Digging deeper variable loadings

Consider the influence of each of the original variables upon the principal components, known as loading scores. We obtain the info using `prcomp()` component `$rotation`. Also summarize with a call to `biplot()`


```
## Lets focus on PC1 as it accounts for > 90% of variance  
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```



Largest positive loading scores “push” N. Ireland to right positive side of the plot because of Fresh_potatoes and Soft_drinks. High negative scores of Fresh_fruit and Alcoholic_drinks push the other countries to the left side of the plot.

Q9: Generate a similar ‘loadings plot’ for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?


```
## focusing on PC2
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```



Negative loading scores include Soft_drinks, Alcoholic_drinks, Processes_potatoes. High positive scores include Fresh_potatoes, Other_veg, and Cereals

Another way to see information together with the main PCA plot is in a biplot

```
## The inbuilt biplot() can be useful for small datasets
biplot(pca)
```



There is a central group of foods in the middle of each PC, with a few that are far from the group. This could be representative of how England, Wales and Scotland were clustered together being “similar”, whilst Northern Ireland was the country that was away from the cluster.

Now we will look at the PCA of RNA-seq data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

```
##      wt1 wt2 wt3 wt4 wt5 ko1 ko2 ko3 ko4 ko5
## gene1 439 458 408 429 420 90 88 86 90 93
## gene2 219 200 204 210 187 427 423 434 433 426
## gene3 1006 989 1030 1017 973 252 237 238 226 210
## gene4 783 792 829 856 760 849 856 835 885 894
## gene5 181 249 204 244 225 277 305 272 270 279
## gene6 460 502 491 491 493 612 594 577 618 638
```

```
View(rna.data)
```

Q10: How many genes and samples are in this data set?


```
dim(rna.data)
```

```
## [1] 100 10
```

There are 100 genes and 10 samples in this data set.


Do a PCA and plot results rather than use any other plots

```
## Again we have to take the transpose of our data
```

```
pca <- prcomp(t(rna.data), scale=TRUE)
```

```
## Simple un polished plot of pc1 and pc2
```

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")
```



Let's examine a summary of how much variation in the original data each PC accounts for

```
summary(pca)
```

```
## Importance of components:
```

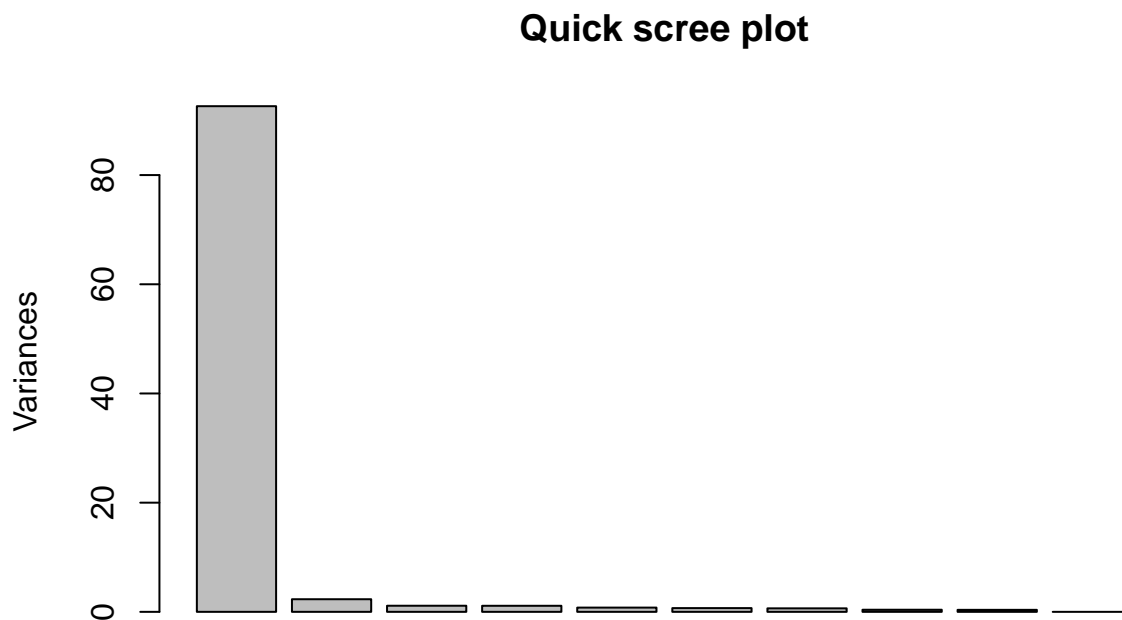
```
##          PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  9.6237 1.5198 1.05787 1.05203 0.88062 0.82545 0.80111
## Proportion of Variance 0.9262 0.0231 0.01119 0.01107 0.00775 0.00681 0.00642
## Cumulative Proportion 0.9262 0.9493 0.96045 0.97152 0.97928 0.98609 0.99251
##          PC8      PC9      PC10
## Standard deviation  0.62065 0.60342 3.348e-15
```

```
## Proportion of Variance 0.00385 0.00364 0.000e+00
## Cumulative Proportion 0.99636 1.00000 1.000e+00
```

By viewing the cumulative proportion, we see that PC1 captures 92.6% of the action. This means that we have successfully reduced a 100 dimensional data set down to only one dimension that retains 92.6% of the essential features from the original data. PC1 captures 92.6%, PC1 and PC2 together captures 94.9%.

Quick barplot summary of this Proportion of Variance for each PC

```
plot(pca, main="Quick scree plot")
```



Square of `pca$sdev` to calculate how much variation in the original data each PC accounts for


```
## Variance captured per PC
pca.var <- pca$sdev^2
```

```
## Percent variance is often more informative to look at
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
pca.var.per
```

```
## [1] 92.6 2.3 1.1 1.1 0.8 0.7 0.6 0.4 0.4 0.0
```

Generate another scree plot

```
barplot(pca.var.per, main="Scree Plot",
        names.arg = paste0("PC", 1:10),
        xlab="Principal Component", ylab="Percent Variation")
```




Making the plot appear more useful

```
## A vector of colors for wt and ko samples
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

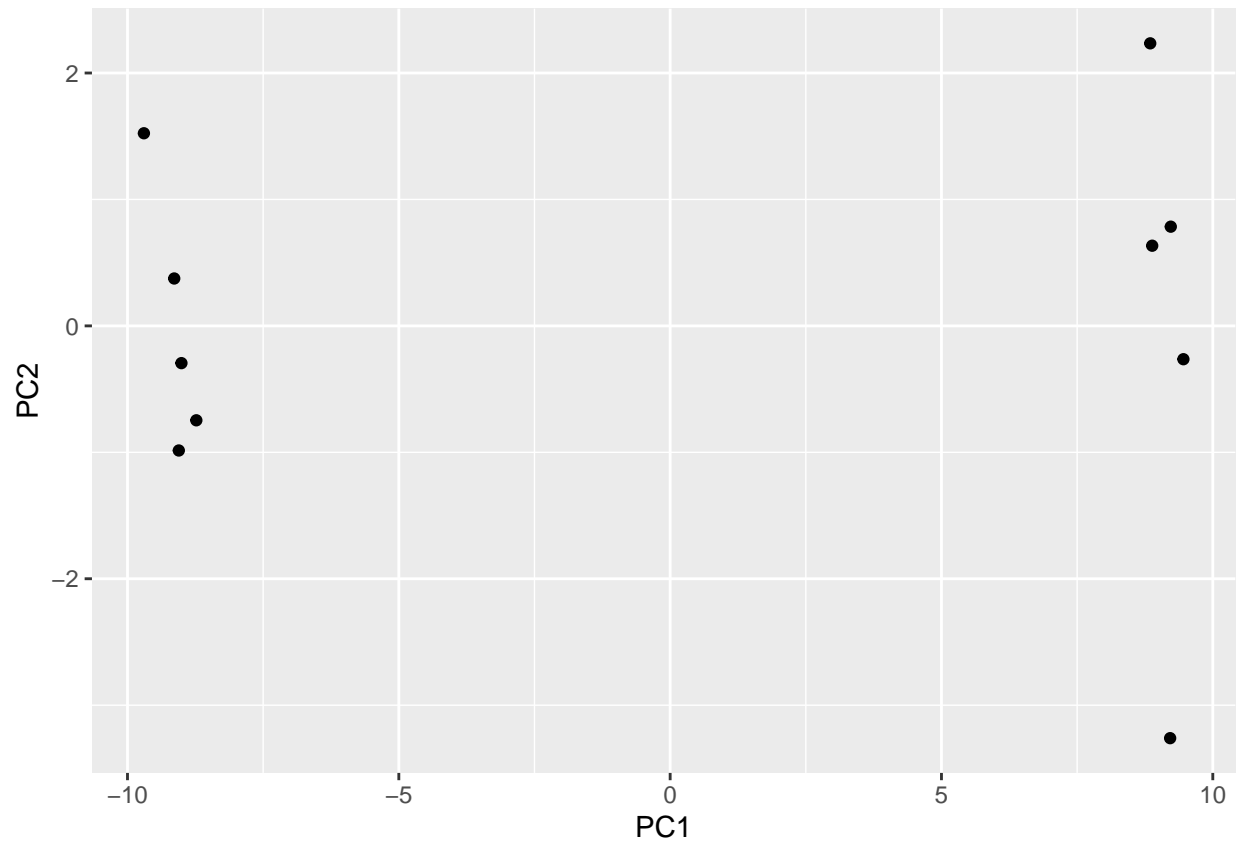
plot(pca$x[,1], pca$x[,2], col=colvec, pch=16,
      xlab=paste0("PC1 (", pca.var.per[1], "%)"),
      ylab=paste0("PC2 (", pca.var.per[2], "%)"))

text(pca$x[,1], pca$x[,2], labels = colnames(rna.data), pos=c(rep(4,5), rep(2,5)))
```



Using ggplot


```
# Prep ggplot
library(ggplot2)
# Contain PCA results in dataframe
df <- as.data.frame(pca$x)
# Our first basic plot
ggplot(df) +
  aes(PC1, PC2) +
  geom_point()
```



Adding some aesthetics and color

```
# Add a 'wt' and 'ko' "condition" column
df$samples <- colnames(rna.data)
df$condition <- substr(colnames(rna.data),1,2)

p <- ggplot(df) +
  aes(PC1, PC2, label=samples, col=condition) +
  geom_label(show.legend = FALSE)
p
```




More polishing to the plot

```
p + labs(title="PCA of RNASeq Data",
  subtitle = "PC1 clearly separates wild-type samples from knock-out samples",
  x=paste0("PC1 (", pca.var.per[1], "%)"),
  y=paste0("PC2 (", pca.var.per[2], "%)"),
  caption="BIMM143 RNASeq example data") +
  theme_bw()
```

PCA of RNASeq Data

PC1 clearly separates wild-type samples from knock-out samples



BIMM143 RNASeq example data

Let's find the top 10 measurements (genes) that contribute most to pc1 in either direction

```
loading_scores <- pca$rotation[,1]

## Find the top 10 measurements (genes) that contribute
## most to PC1 in either direction (+ or -)
gene_scores <- abs(loading_scores)
gene_score_ranked <- sort(gene_scores, decreasing=TRUE)

## show the names of the top 10 genes
top_10_genes <- names(gene_score_ranked[1:10])
top_10_genes
```

```
## [1] "gene100" "gene66" "gene45" "gene68" "gene98" "gene60" "gene21"
## [8] "gene56" "gene10" "gene90"
```